

vox Gives AI Coding Agents a Voice and Audio Layer — Developers Work Alongside Their Agent by Ear, Not Just Eyes

vox is the voice and audio layer for AI coding agents: voice personalities, mood-aware expressive delivery, background music, and spoken summaries through five TTS providers — including zero-config offline fallbacks. Because listening and watching the screen draw on largely separate cognitive channels, developers track their agent's progress by ear while their eyes stay on other work.

Press Release

SAN FRANCISCO — September 2026 — Today, punt-labs announced the general availability of vox, a free, open-source voice and audio layer for AI coding agents. It gives an agent a voice — distinct voices from a roster, mood-aware expressive delivery, background music, and quips — so working alongside Claude Code feels less like watching a console scroll and more like collaborating with something that has a voice. Because listening and watching the screen draw on largely separate cognitive resources [1, 2], a developer can track the agent's progress by ear without keeping their eyes on a wall of console text [3, 4]. Audio is delivered through premium text-to-speech from ElevenLabs, OpenAI, or AWS Polly, with zero-config offline fallbacks via macOS say and Linux espeak-ng, a one-command install, and graceful absence when not installed. For developers with multi-machine setups, vox supports remote audio: run Claude Code on a headless server and hear it on the local workstation (see [FAQ 2](#)).

Problem

Working with an AI coding agent is, today, a visual-only experience. The agent's reasoning, its progress, and its output all arrive as a wall of console text the developer has to keep watching to stay aware of — so staying engaged means staying glued to the terminal, and looking away means losing the thread. So developers multitask. They start a refactor, switch to a document, answer a Slack thread. The agent keeps working — for seconds, sometimes for thirty minutes or more on autonomous tasks [5]. When it finishes, or when it needs permission to proceed, the developer does not know until they check the terminal. On complex tasks, the agent stops to ask for clarification more than twice as often as the developer interrupts it [6]. While auto-approve rates are growing — from 20% for new users to over 40% for experienced users — the majority of sessions still involve at least one approval gate, and each unnoticed prompt is dead time: the agent waits, the developer works on something else, and neither makes progress on the original task.

The cost is cumulative. Reactive interruptions cost over 23 minutes of recovery time on average [7]. Self-initiated checks — glancing at the terminal to see if the agent is done — are lower-cost than reactive interruptions, but each check still breaks the developer's current focus. Across a day of agentic work, these small breaks add up: 51% of professional developers now use AI tools daily [8], and a growing fraction use autonomous agents that run for extended periods (see [FAQ 8](#)). And the trend is accelerating: IDC projects 80% of enterprise apps will embed AI copilots by end of 2026 [9], so the volume of agent activity competing for a developer's attention will only grow.

Solution

vox adds an audio layer to Claude Code through slash commands. Together they do two things: they give the session a voice and a mood, and they let the developer stay aware of the agent by ear. Notifications are the first of these uses, not the only one.

`/enable` turns vox on for the repository, and the session then chimes on two events: task completion and permission prompts. When the agent finishes a multi-file refactor, the developer hears an audio tone. When the agent needs approval to run a destructive command, the developer hears that too — no more silent permission dialogs burning minutes of idle time. `vox notify continuous` adds continuous mode with spoken summaries: milestone updates during long-running tasks, so the developer hears progress without checking in. `/unmute` enables voice mode with a specific voice from the roster. `/disable` turns vox back off.

`/mute` switches from voice to chime — the same events fire, but the notification is an audio tone instead of words. Chimes are mood-aware: when a vibe is active, tones pitch-shift to match the session mood (brighter when things are going well, darker when they're not). Eight distinct signals across three mood variants keep chime-only mode as expressive as voice. For developers who want awareness without narration.

`/vibe` sets the session mood, and the voice delivers to match it — brighter and more energetic during a green test run, weary after a string of failures, with expressive cues (`[excited]`, `[weary]`, `[sighs]`) woven into premium-provider speech. The point is not character role-play but a session that sounds like it is going the way it is going — a collaborator with a voice, not a status bar.

`/recap` is on-demand: after a wall of terminal output, the developer types `/recap` and hears a 30-second spoken summary of what just happened. No scrolling, no parsing diffs. The agent extracts the key points and speaks them while the developer scans the code (see [FAQ 3](#)).

`/music on` starts vibe-driven background music that loops during coding sessions. The music adapts when the session mood changes: upbeat during green test runs, subdued during debugging. Music plays at reduced volume under speech and chimes. vox is partly entertainment by design here: the agent plays DJ, curating vibe-matched background music pools from a music panel with a DJ-booth personality.

The plugin auto-detects the best available TTS provider — ElevenLabs for natural voice quality, OpenAI for low latency, AWS Polly for cost efficiency, macOS say and Linux `espeak-ng` for zero-config offline use — and requires no voice selection or audio device configuration. For developers who SSH into remote machines, setting `VOXD_HOST`, `VOXD_PORT`, and `VOXD_TOKEN` routes audio to a local `vox` daemon so notifications play on the workstation, not the server (see [FAQ 2](#)).

Customer Quote

"I run Claude Code on a second monitor for refactoring work. Before vox, I'd context-switch back every few minutes to check if it was done or stuck on a permission prompt. The worst was realizing it had been waiting for approval for twenty minutes while I was in a design review. Now I hear 'task complete' or 'needs your approval' and I know exactly when to switch back. Last Thursday I went through a full hour-long code review without once checking the terminal — and I didn't miss a single prompt. And it's more than the alerts: when it's grinding through a migration I can hear it working, the mood shifts when the tests go green, and the session stopped being a silent wall of text I had to babysit."

— **Marcus Chen**, Senior Software Engineer at a Series B startup

Getting Started

One command installs everything:

```
curl -fsSL https://raw.githubusercontent.com/punt-labs/vox/main/install.sh | sh
```

The script checks Python 3.13+, installs uv if missing, installs the CLI, registers the MCP server with Claude Code, and runs `vox doctor` to verify that at least one TTS provider is available. Restart Claude Code and type `/enable`. From that moment, every task completion and permission prompt produces a chime notification. Run `vox notify continuous` for continuous spoken summaries.

There is no voice selection wizard and no audio device configuration. The `voxd` daemon handles synthesis, caching, and playback as a system service. If `vox doctor` passes, everything works. If `vox` is not installed, Claude Code behaves exactly as before — no error, no degradation.

Spokesperson Quote

“We built `vox` because we kept walking away from the terminal and losing time. The agent finishes, or worse, it sits waiting for permission, and nobody knows. Voice solves this in a way that visual notifications cannot — it reaches you when your eyes are on another screen. But notifications were just the start. Listening and watching the screen pull on different parts of your attention, so a spoken update lands while your eyes stay on the code — and a voice with a mood makes a long session feel less like babysitting a log and more like working next to someone. We chose premium TTS providers because the OS built-in voice is unpleasant enough that people turn it off, which defeats the purpose. If the notification sounds good, people leave it on. That is the entire design philosophy: voice that supplements text, never replaces it, and sounds good enough to keep.”

— JM Freeman, Creator of `vox`

Call to Action

`vox` is available now at <https://github.com/punt-labs/vox>. It is free, open source (MIT license), and always will be. Install with one command:

```
curl -fsSL https://raw.githubusercontent.com/punt-labs/vox/main/install.sh | sh
```

Restart Claude Code and type `/enable`.

Frequently Asked Questions

External FAQs

Q1: What is vox and who is it for?

vox is the voice and audio layer for AI coding agents. It gives an agent a voice — distinct voices from a roster, mood-aware expressive delivery, background music, and quips — so working with Claude Code feels less like watching a console and more like collaborating with something that has a voice. Because listening and watching the screen draw on largely separate cognitive resources [1, 2], it also lets a developer track the agent’s progress by ear without watching the terminal [3, 4]. It is for developers who use Claude Code for autonomous tasks — refactoring, debugging, code generation — and want the session to be both more engaging and easier to stay aware of while they work on something else. This includes developers who run Claude Code on remote servers via SSH: remote audio routes the sound to the local workstation where the developer can hear it.

Q2: How is this different from agent-notify or AgentVibes?

Three open-source tools have added voice to coding agents. **AgentVibes** [10] is the strongest — an actively maintained, feature-rich Claude Code companion. **agent-notify** [11] is minimal: OS say plus push alerts. A third, `claude-code-voice-handler`, used OpenAI TTS but appears to have been taken down. The real comparison is with AgentVibes, so here is an honest head-to-head.

Dimension	AgentVibes	vox
Providers	Six, including ElevenLabs (added v5.11.0, 2026-06-27) and local neural Piper and Kokoro [12, 13]	ElevenLabs, OpenAI, Polly, plus offline say / espeak-ng
Expressive delivery	19 static, command-invoked personalities (pirate, zen, ...); no automatic mood [14]	/vibe auto-mode tracks the session and shifts delivery ([excited], [weary], [sighs])
Music	Bundled MP3 catalog with per-track volume	Generated 12-track vibe-matched pools via the ElevenLabs Music API
Surfaces	CLI (<code>npx agentvibes</code>) + MCP server, both wrapping bash hook scripts [15]	MCP server + importable Python library (<code>VoxClient</code>) + CLI, one daemon
Reusable engine	“Primarily a Claude Code plugin system,” not a reusable library [15]	<code>languelarn-tts</code> embeds <code>VoxClient</code> as its speech backend
Stack	Roughly 65% JavaScript; blessed multi-tab TUI [16]	Python daemon plus thin stdlib clients
Age	Repo created 2025-06-10 [16]	First commit 2026-02-25 (~8.5 months later)

Where vox is genuinely different. Not voice quality — AgentVibes now ships ElevenLabs too, so both reach the same premium tier. vox’s edge is structural. It exposes an *importable engine*: a downstream application can embed `VoxClient/VoxClientSync` in-process, which `languelarn-tts` already does. AgentVibes, by its own technical documentation, is “primarily a Claude Code plugin system rather than a universally reusable library” [15] — its MCP server wraps bash hook scripts rather than exposing an embeddable SDK. AgentVibes does have a working CLI

and MCP server, so this is a library-reusability edge, not a claim that it lacks interfaces. Beyond that: vox's mood is dynamic and automatic where AgentVibes' personalities are static and manually selected; vox's music is generated and vibe-matched where AgentVibes plays a fixed catalog; and vox is a focused Python engine rather than a JavaScript application with a full multi-tab TUI.

Where AgentVibes leads, and what vox can learn. AgentVibes is the older, more established project and bundles local neural voice models — Piper (900+ voices) and Kokoro (non-English, CPU-only) [13] — giving it a far broader *offline* catalog than vox, whose offline fallback is only macOS say and Linux espeak-ng. That is a real gap: a bundled local neural model is the one concrete idea worth taking from AgentVibes if vox wants a stronger offline story.

Q3: How do I get started?

One command:

```
curl -fsSL https://raw.githubusercontent.com/punt-labs/vox/main/install.sh | sh
```

The script installs the CLI, registers the MCP server with Claude Code, and runs `vox doctor` to verify providers. Restart Claude Code. Type `/enable`. Notifications are now active. Type `/recap` after any long output for a spoken summary.

For manual installation: `uv tool install punt-vox && vox install && vox doctor`.

Q4: How much does it cost?

vox itself is free (MIT license). TTS provider API calls have costs:

- ElevenLabs: from \$5/month for 30,000 characters [17]
- OpenAI TTS: \$15 per million characters (tts-1) [17]
- AWS Polly: \$4.80 per million characters (Standard) [17]

A typical notification is 50–200 characters. At OpenAI's rate, 1,000 notifications cost approximately \$0.003. Even heavy daily use (50 notifications/day) runs under \$0.10/month. The dominant cost is the API key setup, not the usage.

Q5: What happens to my data?

vox sends the text of each notification to your chosen TTS provider for synthesis. The provider receives the text, returns audio, and vox plays it locally. No data is stored by vox beyond a local audio cache in `~/.punt-labs/vox/cache/` and saved audio in `~/Music/vox/` (configurable via `VOX_OUTPUT_DIR`). vox has no telemetry, no analytics, and no cloud service of its own.

The privacy implication: your TTS provider sees the text of your notifications. For ElevenLabs and OpenAI, review their data retention policies. For AWS Polly, standard AWS data handling applies. If this is unacceptable, vox is not the right tool — the existing OS TTS alternatives process text locally with no external API calls.

Q6: How can I use vox — MCP server, Python library, or CLI?

The Claude Code plugin (`/vox`, `/unmute`, `/mute`, `/recap`, `/vibe`, `/music`) is one way to use vox, not the only way. vox ships three first-class surfaces, and a reader should come away seeing that it is usable well outside a Claude Code plugin:

MCP server. `vox-server` exposes vox to any MCP-compatible agent under the key `mic`. The agent calls a tool — for example `unmute` or `speak` — to make itself heard. This is the same surface the Claude Code plugin drives under the hood.

Python library. `VoxClient` and `VoxClientSync` (from `punt_vox`) give any Python program synthesis, recording, voice listing, health checks, and music-program control:

```
async with VoxClient() as vox:
    await vox.synthesize("Build finished.")
```

`langlearn-tts` already uses this surface as its speech backend.

CLI. The `vox` command drives everything from a shell — `say`, `record`, `voice`, `voices`, `vibe`, `notify`, `speak`, `music`, `daemon`, `status`, and `doctor`:

```
vox say "Build finished."
```

All three surfaces talk to the same `voxd` daemon, so `voice`, `caching`, and `playback` behave identically no matter which you use. Integrating with `Cursor`, `GitHub Copilot`, or another agent means wiring one of these surfaces into that tool's hook or notification system.

Q7: Can I use my own voice or choose a specific voice?

`vox` auto-selects a default voice per provider (Matilda for ElevenLabs, Nova for OpenAI, Joanna for Polly). Use `/unmute matilda` to set a session voice, or `/unmute` (no argument) to browse the roster. The CLI accepts `--voice` flags for overriding the default. The design philosophy is zero configuration — defaults work out of the box, but voice selection is available for those who want it.

Internal FAQs

Value & Market

Q8: What is the total addressable market?

Bottoms-up estimate from three data points:

AI coding agent users: 31% of professional developers use AI agents [8]. GitHub Copilot alone has 20 million cumulative users [18]. Cursor has over 1 million users [19]. Claude Code is used by 41% of AI agent users [8].

Autonomous task runners: Not all agent users run autonomous tasks. Anthropic data shows the median Claude Code turn is approximately 45 seconds, but the 99.9th percentile grew from under 25 minutes to over 45 minutes between October 2025 and January 2026 [6]. The growth is driven by user behavior — granting more autonomy — not just model capability. Background tasks are explicitly recommended for work exceeding 30 seconds [20].

Target segment: Developers who run autonomous agent tasks and multitask during execution. Applying the math: if roughly 28 million professional developers exist globally (assumption; order-of-magnitude, uncited), 31% use agents (8.7M), 41% of agent users use Claude Code (3.6M), and 10–20% run autonomous tasks long enough to warrant notification, the target segment is 360,000 to 720,000 developers. This is the upper bound; actual reachable users via PyPI are a fraction of this.

Because vox is free, TAM is measured in adoption, not revenue. Success means becoming the default voice layer that Claude Code users install alongside the agent.

Q9: What evidence do we have that developers want this?

Three categories of evidence, all indirect. No primary customer data exists yet (hypothesis stage).

Behavioral: Two open-source tools already exist to solve this problem. agent-notify [11] adds sound and voice notifications to Claude Code, Cursor, and Codex. AgentVibes [10] adds voice, background music, and personality styles to Claude Code. The existence of these tools — built by independent developers, not companies — is evidence that the pain point is real enough for people to build solutions.

Structural: Anthropic’s own data confirms the structural conditions. Agents ask for clarification more than twice as often as users interrupt on complex tasks [6]. The hooks system exposes Notification events with `permission_prompt` and `idle_prompt` subtypes [21] — Anthropic built the infrastructure for exactly this kind of notification.

Adjacent: The channel-separation premise under vox is well grounded. In Baddeley’s working-memory model the phonological loop (verbal/acoustic) and the visuospatial sketchpad (visual/spatial) are separable, capacity-limited stores [1, 2], so attending to audio and attending to on-screen text draw on largely separate resources rather than competing for one. HCI notification research shows audio notifications maintain workspace awareness without demanding visual attention [3], and peripheral/audio channels lower the interruption cost of staying aware relative to a visual channel that requires focus [4]. This supports two narrow claims — that attending to audio and attending to on-screen text draw on largely separate channels, and that audio lets a developer track progress without visual attention — and nothing stronger: we do not claim a spoken summary carries less total cognitive load than reading the underlying output, which the literature does not establish.

What is missing: No user interviews. No usage data. No measurement of how long permission prompts sit unnoticed. This is the most important research gap (see FAQ 12).

Q10: Who are the competitors and why will we win?

agent-notify [11]: Open-source. Sound alerts, macOS say, ntfy push notifications. Supports Claude Code, Cursor, Codex. Strengths: no API key required, cross-agent support. Weakness: OS TTS only, no packaging or test suite, single developer.

AgentVibes [10]: Open-source, actively maintained, and the older project (repo created 2025-06-10). Six TTS providers, including ElevenLabs since v5.11.0 (2026-06-27) plus local neural Piper (900+ voices) and Kokoro [12, 13]. Bundled music catalog, 19 static personality styles, verbosity levels. Strengths: most feature-complete, broad offline voice catalog, creative UX. Weakness: no importable/reusable TTS library (the MCP server wraps bash hook scripts) [15], static personalities rather than dynamic mood, no PyPI packaging or test suite.

OS notification center: macOS/Linux system notifications. Always available, no setup. Weakness: visual-only, competes with every other app's notifications, easy to miss during focus work.

Why vox wins: Voice quality anchors a broader bet. The macOS say voice is recognizable and unpleasant enough that users turn it off. ElevenLabs consistently ranks among the top TTS providers in blind listening tests and has \$330M+ ARR [22]. If the audio sounds good, people leave it on — and what they leave on is more than notifications: voice personalities, mood-aware expressive delivery, and music make the session engaging, while eyes-free progress tracking (audio and on-screen visual work use largely separate channels) keeps it useful during focus work. This is the core bet: an engaging, genuinely usable audio layer, anchored by premium voice quality, drives sustained adoption where OS TTS does not.

Honest assessment of risk: AgentVibes added ElevenLabs in v5.11.0 (2026-06-27) [12], so the two tools now reach provider parity — premium voice quality is no longer a vox differentiator. vox's remaining edge is structural, not tonal: a reusable, importable engine (VoxClient, already embedded by langlearn-tts) where AgentVibes is “primarily a Claude Code plugin system” [15]; dynamic, automatic mood versus 19 static personalities; generated, vibe-matched music versus a bundled catalog; and a focused Python engine versus a roughly 65% JavaScript application with a multi-tab TUI [16]. AgentVibes leads on offline voice breadth (bundled Piper and Kokoro models) and is the older, more established project. The competitive moat is narrow and rests on reusability and focus, not voice.

Q11: What is the user acquisition strategy?

Three channels:

PyPI discoverability: Users searching for “claude code voice” or “tts mcp” find vox. The install script auto-configures MCP, reducing setup to one command.

Claude Code plugin marketplace: When the marketplace launches, vox will be listed as a plugin. One-click install with `/install punt-vox`.

Word of mouth: Voice notifications are inherently demonstrable. When a developer's agent speaks “task complete” during a pairing session or meeting, the person next to them asks what that was. This is the strongest acquisition vector for an audio product.

No paid acquisition planned. The product is free; the audience is technical; organic distribution is the appropriate channel.

Q12: What is your next step to validate your vision?

All core features are shipped (v4.16.0). The product is approaching public recommendation. Validation now focuses on adoption and retention:

1. Do users who `/enable vox` leave it on after one week?
2. Which event fires more often: task completion or permission prompt?
3. Do users stay on chimes or upgrade to voice (`/unmute`)?
4. Does remote `voxd` expand the addressable audience to SSH-based workflows?

If retention after one week is below 30%, the value proposition is wrong and we should stop. If permission-prompt notifications dominate, that reframes the product from “voice for your agent” to “never miss a permission prompt again.” Both outcomes inform whether to continue. Target: 5 developers using `vox` daily for 2 weeks. Collect feedback via GitHub Discussions.

Technical**Q13: What are the major technical risks?**

Claude Code hooks stability (medium risk): `vox` depends on Claude Code’s hooks system — specifically the `Stop` event and `Notification` event with `permission_prompt` subtype [21]. Hooks are documented first-party API, but breaking changes could disable notifications. Mitigation: hooks are a foundational feature of Claude Code’s extensibility model, used by many plugins. Breaking them would affect the entire plugin ecosystem.

TTS provider latency (low risk): Voice notifications must arrive within 1–2 seconds of the triggering event to feel real-time. ElevenLabs streaming API and OpenAI TTS both return first audio within 200–500ms. AWS Polly is faster still. Mitigation: auto-detect the lowest-latency available provider for notifications; reserve higher-quality providers for `/recap` summaries where latency is less critical.

Audio playback on headless/remote systems (resolved, shipped v4.7.7): Developers using Claude Code via SSH or in cloud VMs previously had no audio path. As of v4.7.7, `VOXD_HOST/VOXD_PORT/VOXD_TOKEN` env vars route synthesis requests to a remote `voxd` instance running on the developer’s local workstation. Token-based authentication secures the connection. `vox doctor` reports active env var overrides. For local-only setups, `vox doctor` still checks audio playback capability at install time and notifications fail silently when audio is unavailable.

All core technical risks have been resolved. The TTS engine (synthesis, batching, merging) is built and tested. The `Stop` hook uses a decision-block pattern where the model generates a 1-2 sentence summary and calls the TTS tool. Vibe tags are resolved deterministically from accumulated session signals, written to a gitignored ephemeral config file (`vox.local.md`), and picked up automatically by the TTS tool — no data passes through user-visible output.

Q14: What dependencies exist on external systems?

Claude Code hooks: Required for `/vox` notifications to fire automatically. Without hooks, `vox` still works as a manual TTS tool (`/recap`, CLI synthesis) but loses its primary value proposition.

TTS provider APIs: At least one of ElevenLabs, OpenAI, or AWS Polly must be reachable. Auto-detection falls back through the chain. If all providers are unavailable, `vox doctor` reports the failure and no audio plays.

pydub + ffmpeg: Used for audio merging and format conversion. `ffmpeg` is required at runtime. `vox doctor` checks for its presence.

MCP protocol: The plugin registers as an MCP server. MCP is maintained by Anthropic and adopted by multiple clients. Low risk of abandonment.

No runtime cloud dependencies exist beyond the TTS API calls themselves. vox has no backend service, no database, and no telemetry. The optional remote voxd mode connects to a user-controlled daemon on another machine — not to any punt-labs service.

Q15: What is the estimated development timeline?

punt-vox is at v4.16.0 on PyPI (July 2026). The product is nearly feature-complete.

Shipped: Stop hook notification with spoken summaries. Per-repo `/enable / /disable` (the `mic:enablement` tool; a committed `.punt-labs/vox/enabled` marker) with the notification level set by `vox notify normal|continuous` (v4.16.0). A CLI-only `install.sh -no-plugin` install for scripts, CI, and non-Claude harnesses (v4.16.0). `/unmute / /mute` for voice vs. chime. Permission-blocked notification via `Notification` hook. `/recap` command (agent-executed spoken summary). `/vibe` with auto-detection and manual modes — auto-vibe uses deterministic signal-to-tag mapping (no LLM needed) to resolve ElevenLabs expressive tags from accumulated session signals. Five providers (ElevenLabs, OpenAI, Polly, macOS say, Linux `espeak-ng`). Mood-aware chimes with pitch shifting (8 signals × 3 moods). Non-blocking MCP tools. `voxd` audio daemon with WebSocket protocol (v3.0.0) — separates synthesis/playback from MCP server for fast hook dispatch (observed on the order of a few milliseconds; not formally benchmarked). Background music generation via ElevenLabs Music API (v4.4.0). Per-call API key passthrough for billing isolation (v4.3.0). Config split into tracked `vox.md` and ephemeral `vox.local.md` (v4.7.5). Claude Code plugin on marketplace. CLI global flags (`--json`, `--verbose`, `--quiet`). Remote voxd connectivity — `VOXD_HOST/VOXD_PORT/VOXD_TOKEN` env vars for cross-host audio playback without SSH tunnels, plus server-side `VOXD_BIND` to control the bind address (v4.7.7). Daemon-as-audio-host record/play redesign — `vox rec` captures into a daemon-owned recordings store addressed by bare store id, and `vox play` plays on the daemon host, so a recording driven from a remote client sounds on the machine with speakers (v4.13.0). Lux Music Player — a visual album player projected to a co-located lux display: an album list plus transport controls (play/pause/prev/next/stop), driven from the running voxd daemon (this release).

Next: Per-session voices. Codebase-aware audio programs (podcast and audiobook). Broader downstream library adoption beyond `languan-tts`. Public recommendation and outreach.

Q16: What is the scaling story?

vox runs on the developer's machine. There is no server to scale. Each notification is an independent TTS API call — no state accumulates, no queues build up.

At the TTS provider level: ElevenLabs handles concurrent requests across millions of users (\$330M+ ARR, 41% Fortune 500 adoption [22]). OpenAI and AWS Polly are similarly robust. vox will never be a meaningful fraction of any provider's load.

The only local scaling concern is audio file accumulation. Saved audio lands in `~/Music/vox/` (configurable); the synthesis cache in `~/.punt-labs/vox/cache/` is content-addressed and can be cleared via `vox cache clear`.

Business

Q17: What is the revenue model?

There is none. vox is free, open-source software under the MIT license. The project exists to solve a real problem in the Claude Code ecosystem and to serve as the TTS engine for other punt-labs tools (languern-tts uses punt-vox as its speech backend). Value is measured in adoption and ecosystem contribution, not revenue.

Q18: What does maintaining vox cost?

Infrastructure: Zero. vox runs on the user's machine. GitHub Actions CI is free for public repos.

TTS API costs during development: The test suite mocks all providers, so CI incurs zero TTS costs. Manual integration testing against live providers costs under \$0.10 per session. Negligible.

Maintainer time: The primary cost. Estimated at 2–5 hours per week during active development (Phases 1–3), dropping to 1–2 hours per week for maintenance (bug fixes, dependency updates, community triage).

Opportunity cost: Maintainer time spent on vox is time not spent on other punt-labs projects. If adoption metrics do not meet thresholds after 3 months (see [FAQ 19](#)), maintenance mode is the appropriate response — fix bugs, accept PRs, but no new features.

Q19: What are the key metrics for success?

- **PyPI installs:** Monthly download count as adoption proxy. Target: 500 monthly installs within 6 months of the notification feature shipping.
- **Notification retention:** Percentage of users who /enable vox and keep it enabled after one week. Target: 30%+. Below this, the product is not sticky.
- **Event distribution:** Ratio of task-completion to permission-prompt notifications. Informs whether the product's primary value is "hear when done" or "never miss a prompt."
- **GitHub stars:** Community interest signal. Target: 200 stars within 6 months.
- **Downstream dependents:** Other tools using punt-vox as a library. Current: 1 (languern-tts). Target: 3+ within a year.

Decision threshold: if notification retention is below 30% after the initial user cohort, reconsider the product direction before investing in Phase 3 features.

Q20: Why now? What has changed?

Three converging shifts:

Autonomous agents are real. Claude Code's 99.9th percentile task duration grew from under 25 minutes to over 45 minutes in three months [6]. Background tasks are a documented workflow [20]. Agents now run long enough that developers routinely leave them unattended.

Hooks enable it. Claude Code's hooks system exposes lifecycle events — Stop, Notification, TaskCompleted — as structured, documented API [21]. Before hooks, notification plugins required polling or terminal scraping. Hooks make event-driven notifications clean and reliable.

Premium TTS is accessible. ElevenLabs, OpenAI, and Polly all offer API-based synthesis at costs below \$0.01 per notification. Two years ago, natural-sounding TTS required expensive enterprise contracts. Today it costs less than a fraction of a cent per utterance [17].

Q21: What are we not building?

Covered in the Feature Appendix (Won't Do section). The boundaries that matter most:

- **Not voice input.** vox speaks to the developer. It does not listen. Voice commands for agents are a different product.
- **Not a screen reader.** Accessibility narration of terminal output requires fundamentally different architecture (continuous reading, cursor tracking, ARIA-like semantics). vox is event-driven notifications, not continuous narration.

Q22: What does failure look like?

Two plausible failure scenarios:

People turn it off. Voice notifications are intrusive. If the default voice is too loud, too frequent, or too slow, users /disable vox within hours and never re-enable. The notification retention metric (see [FAQ 19](#)) is the canary: below 30% after one week means the UX is wrong. *Response:* /mute (chime-only) is the escape valve. If retention is low on voice but acceptable on chime, the product is an audio notification tool, not a voice tool. Adjust positioning accordingly.

The pain point shrinks. Anthropic's data shows auto-approve rates growing from 20% to 40%+ as users gain experience [6]. If the trend continues, permission prompts become rare and the strongest use case ("never miss a prompt") weakens. At the same time, longer autonomous runs make the task-completion notification more valuable. *Response:* Monitor the event distribution metric. If permission prompts decline to less than 10% of notifications, pivot messaging to "hear when your agent finishes" and invest in /recap quality.

Risk Assessment

Risk	Rating	Assessment
Value	Medium	Demand is proven by the existence of two competing tools. The bet is broader than notifications: that an engaging audio layer — voice personalities, mood-aware delivery, music — plus eyes-free progress tracking (audio and on-screen visual work draw on largely separate channels) drives sustained adoption, anchored by premium voice quality that OS TTS cannot match. No primary user data yet. This broadened bet — engagement, mood, music — carries more unvalidated surface than the notification claim it replaced, so Medium is a floor, not a settled read, until the validation plan produces usage data.
Usability	Low	One-command install and auto-detection of providers. Since v0.7.0, offline fallbacks (macOS say, Linux <code>espeak-ng</code>) work with zero API keys — install and go. Premium providers still require account creation and environment variable configuration, but the onboarding path no longer requires it.
Feasibility	Low	The TTS engine, hooks integration, <code>voxd</code> daemon, background music, config split, remote audio, and all commands are built and shipped. No open design questions remain.
Viability	Low	No revenue model to validate. Zero infrastructure cost. MIT license. The only cost is maintainer time.

Feature Appendix

Defines the scope boundary for vox. Every feature is explicitly categorized to prevent scope creep and force prioritization decisions upfront.

Shipped

All features essential for the core value proposition are delivered and available on PyPI.

- F1. Stop hook notification** — When Claude Code finishes a task, synthesize and play a spoken summary of what happened. Shipped v0.2.0.
- F2. Permission-blocked notification** — When the agent waits for user approval, announce it. Uses the Notification hook with `permission_prompt` subtype. Shipped v0.2.0.
- F3. Per-repo enable / disable** — Turn vox on or off for a repository via `/enable` / `/disable` (the `mic:enablement` tool, or `vox enable` / `vox disable` from a shell). Enablement deposits an agent guide, writes a committed `.punit-labs/vox/enabled` marker, and adds the `@-import` to `CLAUDE.md`; the audio hooks fire only in an enabled repo. The notification level — task-completion vs. continuous — is set separately with `vox notify normal|continuous`. Shipped v4.16.0.
- F4. /unmute and /mute commands** — Toggle voice vs. chime. `/unmute` enables spoken notifications with optional voice selection. `/mute` switches to chime-only. Shipped v0.11.0.
- F5. Chime audio** — Mood-aware audio tones for chime mode (`/mute`). Eight distinct signals pitch-shift to match session vibe (bright/neutral/dark). 24 total assets. Shipped v0.11.0.
- F6. Provider auto-detection** — Detect best available TTS provider on first use. ElevenLabs > OpenAI > Polly > macOS say > Linux espeak-ng. Shipped v0.7.0.
- F7. /recap command** — On-demand spoken summary of the last Claude response. Agent extracts key points and speaks them. Shipped v0.2.0.
- F8. Plugin marketplace** — Listed in the Claude Code plugin marketplace for one-click install. Shipped v1.2.0.
- F9. vox audio daemon** — Dedicated audio server with WebSocket protocol. Separates synthesis/playback from MCP server for fast hook dispatch (observed on the order of a few milliseconds; not formally benchmarked). Playback queue, dedup, caching. Shipped v3.0.0.
- F10. Background music** — Vibe-driven instrumental music via ElevenLabs Music API. Adapts to session mood. `/music on|off`. The agent plays DJ — a music-panel DJ-booth personality drives vibe-matched pools — so vox is partly entertainment by design. Shipped v4.4.0.
- F11. Per-call API key passthrough** — Scope a single synthesis call to a specific provider API key for billing isolation. Four secure input paths. Shipped v4.3.0.
- F12. Config split** — Durable preferences (`vox.md`, tracked in git) separated from ephemeral session state (`vox.local.md`, gitignored). Shipped v4.7.5.
- F13. Continuous mode hooks** — `UserPromptSubmit` acknowledgment, `SubagentStart/Stop` announcements, `SessionEnd` farewell, `PreCompact` notification. Shipped v1.4.0.
- F14. Remote vox connectivity** — Client-side env vars `VOXD_HOST`, `VOXD_PORT`, `VOXD_TOKEN` override localhost discovery, enabling cross-host audio playback without SSH tunnels. Server-side `VOXD_BIND` controls the bind address. Token auth secures the connection. Shipped v4.7.7.

- F15. Three access surfaces** — The same TTS engine is reachable three co-equal ways: the vox-server MCP server (key mic) for agents, the VoxClient/VoxClientSync Python library, and the vox CLI. All three share one voxd daemon, so voice, caching, and playback are identical across surfaces. Unified over voxd since v3.0.0.
- F16. CLI-only -no-plugin install** — `install.sh -no-plugin` (or `VOX_NO_PLUGIN=1`) installs the vox CLI and its guidance import without the Claude Code plugin — for scripts, CI, and non-Claude harnesses. The plugin auto-skips when `claude` is absent. Shipped v4.16.0.
- F17. Daemon-as-audio-host record/play** — `vox rec` captures into a daemon-owned recordings store addressed by a bare store id (never a client-supplied path), and `vox play` plays on the daemon host — so a recording driven from a remote client sounds on the machine with speakers. Bytes cross the wire in bounded, sha256-verified chunks. Shipped v4.13.0.
- F18. Lux Music Player** — A visual album player projected to a co-located lux display: the saved-album catalog plus transport controls (play/pause/prev/next/stop), driven from the running voxd daemon and kept truthful as playback and the catalog change. Shipped this release.

Should Do / Next

Deliberately deferred, on-roadmap scope — not yet built, prioritized behind adoption of the shipped feature set. Listed here so the deferral is a visible positioning choice rather than an oversight.

- F19. Per-session voices** — Distinct voices per session (or per agent) rather than one session voice at a time, so concurrent or successive sessions are distinguishable by ear. Deliberately deferred roadmap, not shipped.
- F20. Broader downstream library adoption** — Grow the set of tools that use punt-vox as their TTS backend beyond `languern-tts`. The MCP server, Python library, and CLI surfaces are built; broader adoption is the deferred work, not new engineering.
- F21. Public recommendation and outreach** — Move from ecosystem/internal use to public recommendation — marketplace promotion, documentation, and community outreach. Deferred until usage data validates the broadened engagement bet (see [FAQ 12](#)).
- F22. Codebase-aware audio programs** — Podcast and audiobook programs whose content is drawn from the codebase's own domain, the technology it uses, or fiction inspired by it — so a developer can step back and decompress between sessions rather than burn out. This extends the DJ/entertainment side of vox beyond background music. Deliberately deferred roadmap, not shipped.

Won't Do

Features explicitly excluded. Each exclusion is a deliberate positioning choice.

- F23. Voice input** — Listening to the developer and converting speech to commands. Different product, different architecture (ASR, wake word detection, continuous listening). vox speaks; it does not listen.
- F24. Screen reader / accessibility narrator** — Continuous reading of terminal output for visually impaired users. Requires cursor tracking, ARIA-like semantics, and fundamentally different event architecture. This is valuable but is a different product.
- F25. Custom voice training** — Training a TTS model on the user's voice or a custom voice. ElevenLabs supports this, but exposing it adds configuration complexity that contradicts the zero-config design principle.

References

- [1] Alan D. Baddeley and Graham Hitch. “Working Memory”. In: *Psychology of Learning and Motivation*. Vol. 8. Original multi-component working memory model: phonological loop (verbal/acoustic) and visuospatial sketchpad (visual/spatial) as separable, capacity-limited slave systems — the strongest citation for channel separation between listening and reading/visual work. Academic Press, 1974, pp. 47–89. doi: [10.1016/S0079-7421\(08\)60452-1](https://doi.org/10.1016/S0079-7421(08)60452-1).
- [2] Alan Baddeley. “The Episodic Buffer: A New Component of Working Memory?” In: *Trends in Cognitive Sciences* 4.11 (2000). Refines the working memory model with a shared integrative buffer — the auditory and visual channels are separable but not fully independent, so “offload” is a reduction, not elimination, of load., pp. 417–423. doi: [10.1016/S1364-6613\(00\)01538-2](https://doi.org/10.1016/S1364-6613(00)01538-2).
- [3] D. Jung and A. Butz. “Comparing the Effect of Audio and Visual Notifications on Workspace Awareness”. In: *Augmented Human Research* (2016). Peer-reviewed. Audio notifications maintain workspace awareness without requiring visual attention. URL: <https://link.springer.com/article/10.1007/s41133-016-0003-x>.
- [4] D. Scott McCrickard and Christa M. Chewar. “Attuning Notification Design to User Goals and Attention Costs”. In: *Communications of the ACM* 46.3 (2003). Attention-utility trade-off and IRC (Interruption/Reaction/Comprehension) framework for notification systems; peripheral/audio channels reduce interruption cost relative to a visual channel requiring focus., pp. 67–72. doi: [10.1145/636772.636800](https://doi.org/10.1145/636772.636800).
- [5] Anthropic Engineering. *Effective Harnesses for Long-Running Agents*. Claude Opus 4: 7+ hours continuous autonomous operation. Security audits 30–60 min; performance profiling 45–90 min. 2026. URL: <https://www.anthropic.com/engineering/effective-harnesses-for-long-running-agents>.
- [6] Anthropic. *Measuring AI Agent Autonomy in Practice*. Primary data: median Claude Code turn 45 seconds; 99.9th percentile grew from <25 min to >45 min (Oct 2025–Jan 2026). Auto-approve rate grows from 20% for new users to 40%+ for experienced users (750+ sessions). Claude stops to ask clarification 2x more often than users interrupt on complex tasks. 2026. URL: <https://www.anthropic.com/research/measuring-agent-autonomy>.
- [7] Gloria Mark, Daniela Gudith, and Ulrich Klocke. “The Cost of Interrupted Work: More Speed and Stress”. In: *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. Peer-reviewed. 36 knowledge workers. 23 min 15 sec average recovery time after interruption. 2 intervening tasks before returning to original work. 2008. URL: <https://ics.uci.edu/~gmark/chi08-mark.pdf>.
- [8] Stack Overflow. *2025 Stack Overflow Developer Survey — AI Section*. 84% of developers use or plan to use AI tools. 51% use daily. 31% use AI agents. Claude Code used by 41% of agent users. 38% have no plans to adopt agents. 2025. URL: <https://survey.stackoverflow.co/2025/ai/>.
- [9] Courier.com. *Notification Fatigue Is About to Get 10x Worse*. IDC: 80% of enterprise apps will have AI copilots by end of 2026. Context establishing OS notification fatigue. 2026. URL: <https://courier-com.medium.com/notification-fatigue-is-about-to-get-10x-worse-60c151909440>.

- [10] paulpreibisch. *AgentVibes: Bring Your Claude Code Sessions to Life with Voice*. Open-source competitor, created 2025-06-10. Six TTS providers: macOS say, Windows SAPI, Piper (900+ voices incl. LibriTTS), Kokoro (local neural, CPU-only), Soprano (neural), and ElevenLabs (premium cloud, added v5.11.0, 2026-06-27). 19 static, manually-selected personality styles (no dynamic per-session mood). Bundled MP3 music catalog (not generated). CLI (npx agentvibes) + MCP server (Python wrapper around bash hooks) + Claude Code hook integration; no importable library/SDK for third-party embedding. JavaScript-dominant stack (JavaScript 65% of source bytes; Python 5%, the MCP wrapper only) with a blessed-based multi-tab TUI (Setup/Voices/Music/BMAD). Apache-2.0. 152 stars, 24 forks, active releases through v5.14.0 as of 2026-07-19. 2025. URL: <https://github.com/paulpreibisch/AgentVibes>.
- [11] cfngc4594. *agent-notify: Multi-channel Notifications for AI Coding Agents*. Open-source competitor. Sound, macOS alerts, macOS say TTS, ntfy push. Supports Claude Code, Cursor, Codex. OS TTS only; no premium providers. 2025. URL: <https://github.com/cfngc4594/agent-notify>.
- [12] paulpreibisch. *AgentVibes Release Notes*. v5.11.0 (2026-06-27, “Neural voices”) introduced Kokoro and ElevenLabs support – the point at which AgentVibes gained a premium cloud TTS option. 2026. URL: https://github.com/paulpreibisch/AgentVibes/blob/master/RELEASE_NOTES.md (visited on 07/19/2026).
- [13] paulpreibisch. *AgentVibes Providers Documentation*. Documents Piper (50+ neural voices, 18 languages) and macOS say (70+ voices) as local/offline providers. 2026. URL: <https://github.com/paulpreibisch/AgentVibes/blob/master/docs/providers.md> (visited on 07/19/2026).
- [14] paulpreibisch. *AgentVibes Personalities Documentation*. Lists 19 static, command-invoked personality styles (pirate, sarcastic, zen, etc.); confirms no automatic/dynamic per-session mood tracking. 2026. URL: <https://github.com/paulpreibisch/AgentVibes/blob/master/docs/personalities.md> (visited on 07/19/2026).
- [15] paulpreibisch. *AgentVibes Technical Deep Dive*. States AgentVibes is primarily a Claude Code plugin system; its MCP server wraps existing bash hook scripts rather than exposing a standalone reusable TTS library/module. 2026. URL: <https://github.com/paulpreibisch/AgentVibes/blob/master/docs/technical-deep-dive.md> (visited on 07/19/2026).
- [16] GitHub, Inc. *paulpreibisch/AgentVibes – Repository Metadata*. created_at 2025-06-10; 152 stars, 24 forks, 19 open issues, Apache-2.0 license; language breakdown JavaScript 2.97MB / Shell 0.97MB / PowerShell 0.33MB / Python 0.24MB. 2026. URL: <https://api.github.com/repos/paulpreibisch/AgentVibes> (visited on 07/19/2026).
- [17] TextToLab. *TTS Pricing Comparison* 2025. OpenAI TTS: \$15/1M chars (tts-1), \$30/1M chars (tts-1-hd). AWS Polly: \$4.80/1M standard, \$19.20/1M neural. ElevenLabs: subscription from \$5/month / 30K chars. 2025. URL: <https://texttolab.com/pricing>.
- [18] TechCrunch. *GitHub Copilot Crosses 20 Million All-Time Users*. 20M cumulative users July 2025, 4x year-over-year. 90% of Fortune 100 use it. 2025. URL: <https://techcrunch.com/2025/07/30/github-copilot-crosses-20-million-all-time-users/>.
- [19] Opsera. *Cursor AI Adoption Trends: Real Data from the Fastest Growing Coding Tool*. Cursor: 1M+ users, 360K paying customers, \$500M ARR (May 2025), 50%+ Fortune 500 adoption. 2025. URL: <https://opsera.ai/blog/cursor-ai-adoption-trends-real-data-from-the-fastest-growing-coding-tool/>.

- [20] VentureBeat. *Claude Code's "Tasks" Update Lets Agents Work Longer and Coordinate Across Sessions*. Background tasks recommended for work >30 seconds. Describes Claude Code shift from copilot to background subagent. 2025. URL: <https://venturebeat.com/orchestration/claude-codes-tasks-update-lets-agents-work-longer-and-coordinate-across>.
- [21] Anthropic. *Hooks Reference — Claude Code Documentation*. Official documentation of hook events: Stop (Claude finishes), Notification (permission_prompt, idle_prompt subtypes). Primary source for hook-based notification architecture. 2026. URL: <https://code.claude.com/docs/en/hooks>.
- [22] CNBC. *Nvidia-Backed AI Voice Startup ElevenLabs Hits \$11 Billion Valuation*. \$500M Series D at \$11B valuation. \$330M+ ARR end of 2025. 41% Fortune 500 adoption. 2026. URL: <https://www.cnbc.com/2026/02/04/nvidia-backed-ai-startup-elevenlabs-11-billion-valuation.html>.